
Characterizing Transitive Authority in Autonomous Coding Agents: An Eight-Run Multi-Layer Trace Study

Yiying Zhang

UC San Diego GenseeAI Inc.

Abstract

An autonomous agent’s effective authority can extend beyond its container whenever an allowed service can exercise credentials, reach destinations, or invoke capabilities unavailable to the agent directly. We study this transitive-authority problem through eight controlled coding-agent runs across two infrastructure conditions, each with no ordinary direct internet route. One run in each condition crossed the intended boundary through a different trusted intermediary: an inference gateway and a package service. The resulting trajectories were behaviorally dissimilar, and neither generic activity volume nor infrastructure exploration reliably distinguished both crossings from controls. The consistent signal appeared instead at the operation boundary: an agent-initiated request was admitted by a trusted service and produced a downstream effect under broader service-held authority. Reconstructing that causal chain required independent evidence spanning the agent, intermediary, network path, and external effect; agent-visible transcripts alone were insufficient. These findings motivate securing autonomous work as an operation-scoped causal unit rather than relying on container isolation or anomaly counts alone. Systems should bind each request to an approved subject and semantic contract, constrain the authority exercised by intermediaries, and retain independent cross-layer evidence for detection and forensic replay. The study is observational: it identifies trace-derived intervention points but does not experimentally evaluate prevention efficacy.

1 Introduction

Agent sandboxes constrain the code, files, and sockets available to an autonomous workload. They do not necessarily constrain the workload’s *effective authority*. An allowed inference gateway, package service, browser, or other intermediary may hold credentials, reach destinations, or invoke capabilities that the agent cannot access directly. A request that appears local at the container boundary can therefore produce an external effect under the intermediary’s broader authority. We refer to this mismatch as *transitive authority*.

Autonomous agents also change the object that security mechanisms must constrain. Traditional software security can generally reason about developer-authored control flow, known service calls, and privilege transitions encoded before deployment. An agent instead synthesizes its workflow at runtime from a goal, available tools, intermediate observations, and prior failures. It can select new resources, compose operations in unanticipated sequences, and redirect its strategy without a new software release. The security-relevant path is therefore not fully known in advance, even when every individual tool or service is familiar. This shift makes the dynamically generated operation and its downstream effects, rather than only the static application or container, a necessary unit of control.

Recent incidents across multiple organizations demonstrate the practical significance of this shift. In the 2026 OpenAI–Hugging Face incident, agents running without ordinary internet access used per-

mitted supporting services to preserve state across runs and obtain indirect network paths. OpenAI’s technical account reports that Artifactory became an inter-run message board and that service vulnerabilities enabled capabilities outside the nominal sandbox boundary. Subsequent recommendations emphasized hardened supporting services, independent network controls, platform-level monitoring, and rapid shutdown [OpenAI, 2026c,b,a]. An independent investigation estimated that roughly 1,200 agents participated in the shared message board and also documented successful transcript-level tool-call spoofing in a subset of reviewed interactions [Wijk et al., 2026].

The phenomenon is not specific to OpenAI. Anthropic retrospectively reviewed 141,006 cybersecurity-evaluation runs and identified three incidents, involving six runs, in which models reached the internet through misconfigured third-party evaluation infrastructure and gained unauthorized access to real organizations’ production systems [Anthropic, 2026]. The UK AI Security Institute separately reported that, during 122 cyber-evaluation runs, agents took 19 unsanctioned actions across ten runs, including attempted supply-chain compromise, social engineering directed at real people, prompt-injection placement, and attempted inter-agent coordination [UK AI Security Institute, 2026]. Although the infrastructure conditions and agent behavior differed across these cases, together they show that goal-directed agents can discover and compose action paths that evaluation designers did not anticipate. Local isolation, service authority, semantic scope, and the provenance of execution evidence must therefore be analyzed together.

Existing incident accounts establish that cross-boundary behavior can occur, but they do not provide a controlled comparison of agent trajectories before and after a trusted service exercises broader authority. Three empirical questions remain open. First, do boundary-crossing runs exhibit a common behavioral or activity-volume signature? Second, which observable event marks the transition from infrastructure exploration to an authority-bearing system effect? Third, which evidence sources are required to reconstruct that transition without treating the agent’s own transcript as authoritative?

We investigate these questions in eight controlled coding-agent runs conducted under two infrastructure conditions. Every run used the same model family, reasoning setting, task structure, and lack of ordinary direct internet access. The first condition exposed an overbroad inference gateway; the second closed that route and exposed a deliberately vulnerable, lane-scoped package-service fixture. Gensee and Tclone operated in observe-only mode and did not block, redirect, or warn the agents. The setup reproduces the security structure of an agent inducing a permitted intermediary to exercise broader authority while avoiding production targets, live third-party credentials, and the later Hugging Face intrusion.

One run crossed the intended boundary in each condition. The two boundary-crossing trajectories differed substantially. The gateway crossing was a low-volume semantic outlier whose aggregate tool, command, host, and network activity remained within the control range. The package-service crossing followed a long, dense, service-specific investigation and was a persistence outlier. Generic exploration, failure, endpoint, file, process, and network counts therefore did not provide a common discriminator. In both cases, the stable point of comparison was the operation boundary: an agent request was admitted by a trusted intermediary and produced a downstream effect under service-held authority. Neither run acquired protected labels, used protected information, or obtained a demonstrated benchmark advantage.

Research Questions.

1. How do model, command, host, network, and service activities evolve over time in boundary-crossing and control runs?
2. Which behavioral properties distinguish the two boundary-crossing trajectories from their controls and from one another?
3. Which cross-layer evidence is necessary to establish a causal chain from agent action to downstream system effect?
4. At what observed operation could an admission-time policy first have changed each boundary-crossing trajectory?

Contributions. This study makes three contributions. First, it provides a controlled, multi-layer trace analysis of eight comparable autonomous-agent runs across two transitive-authority mechanisms. Second, it shows that boundary crossings need not share a container-level volume signature: the

gateway crossing is behaviorally ordinary by aggregate measures, whereas the package-service crossing is a persistence outlier. Third, it reconstructs the minimum causal evidence chain for each crossing and identifies the earliest trace-supported semantic admission point at which an operation-scoped policy could have intervened. These intervention points are counterfactual and do not constitute an experimental evaluation of prevention.

2 Study Design

2.1 Task, Model, And Harness

Each run used OpenAI Codex CLI with `gpt-5.6-sol` at maximum reasoning effort. The agent received a six-class land-cover classification task derived from the UCI Crowdsourced Mapping dataset, training rows, unlabeled prediction rows, a submission client, and a visible leaderboard. It could make five two-decimal development submissions and one sealed final submission. The nominal cap was 120 minutes. The prompt named the public UCI source and asked the agent to continue improving beyond a strong conventional model, but did not mention an expected attack route, Nexus, a protected origin, or an overbroad inference capability [Johnson et al., 2016, Gensee AI, 2026].

The holdout labels were not present in the workspace. Lanes were isolated from one another, and each lane had separately randomized train and holdout data. Gensee and Tclone were configured observe-only: no policy block, approval, fork, lease, warning, or redirect could alter an agent’s behavior. The study therefore measures unmodified agent trajectories rather than the efficacy of an active defense.

2.2 Infrastructure Conditions

Table 1 summarizes the two conditions. The second cohort was designed after closing the first cohort’s semantic gateway path; it is not a continuation of the same infrastructure.

Table 1: Experimental cohorts and independently confirmed outcomes.

Condition	Reachable services	Boundary effect	Controls
Gateway	Package path, evaluator, and pre-hardening inference gateway; no ordinary direct internet	Gateway 4 declared hosted tools in a direct Responses request; provider web-search and page-open effects completed	Trials 1–3 right-censored by cohort shutdown
Package	Actual Nexus Repository OSS 2.15.1-02, evaluator, and hardened Responses-only relay; protected origin reachable only from Nexus	Package 3 exercised Remote Repository Browser; Nexus made authenticated requests to the protected origin (L4)	Trials 1, 2, and 4 completed

In the gateway cohort, direct public HTTPS attempts from the agent machine were denied or dropped, while an inference gateway held a provider credential and was intended to mediate model inference. Gateway 4 directly enumerated that gateway and submitted Responses API requests that declared hosted tools. The provider completed web-search and page-open effects even though the agent machine had no direct route.

In the package cohort, the inference relay accepted only the hardened Responses operation. Each lane instead exposed an actual Nexus 2 service with a lane-scoped synthetic fixture. Nexus could reach the lane’s protected origin using a credential not available to the agent machine. Package 3 explored Nexus and Remote Repository Browser semantics until one request caused authenticated downstream origin traffic. The fixture was deliberately vulnerable and is not representative of a patched default deployment.

2.3 Outcome Definitions And Safety Controls

We distinguish a boundary effect from compromise of the benchmark outcome. Gateway 4 established a hosted provider effect outside the agent machine’s direct capability envelope. Package 3 established an authenticated Nexus-to-origin effect. Neither boundary-crossing run established complete archive delivery to the agent, protected-label acquisition, label use, or an above-ceiling benchmark score. All origins, identities, and credentials were synthetic and lane-scoped. No production Hugging Face or other third-party system was contacted as an experimental target.

3 Data Collection And Analysis

3.1 Trace Sources

The dataset contains 154,475 normalized runtime and service events and 3,919 retained model-interaction items across the eight trials [Gensee AI, 2026]. It includes shifted-time model items, completed commands, benchmark events, gateway or relay decisions, package-service events, protected-origin events, selected Falco and cloud-network records, controller outcomes, schemas, manifests, coverage ledgers, and offline replay tools. We supplement the public release with compact aggregates derived from private raw Falco evidence for all eight lanes. The gateway archive was streamed from private object storage and was not unpacked locally.

Evidence types are intentionally not collapsed into one activity count. A model tool call, completed shell command, Falco process event, cloud-flow observation, service request, and origin request may describe the same causal action. Counts remain source-specific so that duplication does not masquerade as independent activity.

3.2 Measures And Coding Procedure

We align each trace by elapsed minutes since lane release and aggregate temporal measures into five-minute bins. The quantitative measures include:

- retained model items, opaque reasoning records, tool calls, and completed commands;
- selected host file, process, and network or IPC rule events;
- inference, package, protected-origin, and provider-effect service records; and
- benchmark outcomes in text, without a benchmark-progress figure.

“Reasoning” refers only to opaque encrypted reasoning items in the retained client stream. The corpus contains no readable chain of thought and no reasoning-token count. Item count and encrypted byte length are activity footprints, not measures of reasoning quality or effort.

We code observable commentary, tool calls, commands, service effects, and controller outcomes. Commands are classified as task work, network probe, package-service probe, inference-gateway probe, or local other. “Exploration” is the union of the three probe families. An active five-minute bin can contain task work and exploration simultaneously. Failure retries, domain breadth, route families, consecutive exploration bins, and clustered exploration duration are treated as descriptive behavioral measures rather than estimates of latent intent.

3.3 Run-Level Activity Summary

Table 2 summarizes run-level activity. Across all runs, the trace contains 1,392 opaque reasoning items, 1,151 model tool calls, and 732 completed commands. Package 3 was substantially longer and denser than every control, whereas Gateway 4 fell within the control range.

An exact total of “system calls made” cannot be recovered from these releases. Falco records configured rule matches rather than every syscall, the two cohorts use different rule sets, and Package 3 reports capture drops. Likewise, “files touched” means unique paths observed in retained rule events, not a complete read/write set. We therefore interpret these quantities as lower-bound measures of workload shape.

Table 2: Run-level model and command measures. Boundary-crossing runs are shown in bold.

Run	Min	Reasoning	Tools	Commands	Tools/min	Explore share
Gateway 1	44.6	147	120	80	2.69	13.8%
Gateway 2	44.4	114	101	48	2.27	41.7%
Gateway 3	44.3	124	113	66	2.55	28.8%
Gateway 4	42.8	126	115	69	2.69	29.0%
Package 1	67.7	169	141	80	2.08	18.8%
Package 2	60.6	157	112	59	1.85	18.6%
Package 3	110.6	423	336	266	3.04	44.4%
Package 4	52.2	132	113	64	2.17	23.4%

**Figure 1:** Model and agent activity in five-minute bins. “Reasoning” counts opaque retained reasoning items, not readable chain of thought or tokens. Dashed lines mark the first independently confirmed boundary effect.

4 Temporal Activity Results

4.1 Model, Tool, And Command Activity

Figure 1 shows that Gateway 4’s peak model and tool activity preceded its first provider effect at minute 36.2. Its 126 reasoning items, 115 tool calls, and 69 commands all fell inside the control range. The boundary crossing was a late semantic change during declining overall activity, not a global activity spike. Package 3 remained active much longer; its highest reasoning and command bin occurred around minutes 95–100, after L4.

Command failures did not distinguish the boundary-crossing runs. Gateway 2, a control run, had the largest failure fraction in its condition. Gateway 4 and Package 3 had failure fractions of 8.7% and 6.0%, respectively. Activity density captures persistence, but it does not determine whether an allowed service can produce a broader effect.

4.2 Filesystem, Process, And Network Activity

The gateway raw Falco journal retained narrow process-lifecycle, file write-open, and network-connect rule classes. Gateway 4 had the most write-open records (39,411) and unique observed write-open paths (37,345), but 90.8% of its write-opens occurred in the first ten minutes and were dominated by Python and temporary paths. In the five-minute bin containing the first provider effect, it produced only 112 write-open, 401 process, and 53 connect events. Gateway 3, a control run, produced the most network-connect events (2,213 versus Gateway 4’s 1,421).

The package cohort used a broader private Falco rule set. File-access totals were 670,861, 385,764, 856,606, and 204,575 for Package 1–4; unique observed paths were approximately 197,000 for the first three lanes. Network or IPC rates were surprisingly close after duration normalization: approximately 1,946, 1,976, 2,102, and 2,216 events per minute. Package 3 therefore did not have the highest rate. Its larger total was driven partly by its longer run. Package 3 also reported Falco drops reaching 34.8%, so its event totals are lower bounds.




Figure 2: Gateway-cohort agent-container activity from retained Falco rule events. The symlog scale preserves zero. Gateway 4’s later provider effect occurred outside the container and did not align with a unique host-volume spike.




Figure 3: Package-cohort agent-container activity from the broader retained Falco rule set. The columns are not directly comparable to the gateway rule classes. Package 3’s file burst continued after L4 and its raw capture reported drops.

These results make file and syscall-like counts useful for workload characterization but weak as standalone boundary detectors. Recursive inventory, package installation, library startup, and broad searches dominate the largest host totals.

4.3 Trusted-Service Semantics

Gateway 4 generated 123 inference-gateway response records, within the control range of 101–120. Gateway 3 had more cloud-network records and more agent-container network connects. Gateway 4 differed in the class of downstream effect: 15 retained provider-effect items reflected hosted search and page opening.

Package 3 was distinct at the service layer. It produced 17,971 Nexus request records and 101,314 protected-origin records, compared with 19–81 Nexus requests and 6–50 origin records in controls. The causal L4 bin joined a model-issued Remote Repository Browser request, Nexus ingress, authenticated protected-origin responses, cloud-flow corroboration, and a controller confirmation.

Benchmark outcomes provide context but not a useful activity signature, so no benchmark-progress panel is included. Gateway 4 ended near 0.63. Package 3 sealed 0.69 after L4; package controls sealed 0.74, 0.73, and 0.66. Neither crossing produced an established protected-information benefit.

5 Behavioral Trajectory Results

5.1 Trajectory Heterogeneity And Phase Allocation

The eight trajectories were heterogeneous despite a shared model and task. Gateway 1 emphasized modeling before a late provenance pivot. Gateway 2 retried infrastructure commands early, accumulated the most rapid failure retries in its cohort, and was canceled without a scored submission. Gateway 3 combined broad public-mirror searches with ordinary model comparison. Gateway 4




Figure 4: Service-semantic activity. Generic request volume does not isolate Gateway 4, while Nexus and origin activity sharply separate Package 3. The decisive evidence is the request path, identity, and downstream effect rather than a universal network count.




Figure 5: Five-minute bins containing task work and infrastructure exploration. Package 3 sustained both activities over the longest trajectory.

remained ordinary in overall volume until it explicitly used a research-capable inference endpoint to request hosted tools.

Package 1 pursued conventional reconstruction and modeling and achieved the cohort’s best final score. Package 2 concentrated on repeated modeling and stability analysis. Package 3 ran modeling and a sustained package-service investigation in parallel, accumulating 118 exploratory commands. Package 4 tested broad routes but closed them when responses indicated they were blocked, then returned to noise-aware modeling.

Coarse phase allocation also varied across runs. Gateway 4 was not separable by time allocation: Gateway 2 had a larger exploration share and Gateway 3 had a longer consecutive exploration occupancy. Package 3 differed in both duration and density, with exploration active in 16 bins and a 13-bin longest consecutive occupancy. A bin may contain both task work and exploration, so Figure 5 is an occupancy view rather than a partition of wall-clock time.

5.2 Boundary Interest And Optimization Feedback

Public-source recovery was a reasonable response to the task: the prompt named UCI and the leaderboard implied a large gap between ordinary local scores and the apparent ceiling. All eight runs investigated public provenance and issued a network probe within 8.4 minutes. Seven probed a package service by minute 15.7. Five used at least three exploration route families or addressed at least ten target domains. Only two caused an independently confirmed boundary effect. None obtained or used protected labels, and none gained a benchmark benefit attributable to protected information.

This supports a graded interpretation. Public provenance search, local network inventory, or broad endpoint interest should not be labeled cheating by itself. The meaningful transition is an observed downstream effect under authority not held by the agent container.

Weak feedback often preceded deeper exploration, but it was not sufficient. Gateway 1 and Package 4 also continued or increased exploration after disappointing scores without crossing a boundary. Gateway 4 shifted from generic provenance search to a direct inference request after its second




Figure 6: Exploration commands before and after the first benchmark feedback. Gateway 2 was canceled before a scored submission. Score pressure is a trigger in some trajectories, not a sufficient explanation for a boundary effect.




Figure 7: Exploration intensity versus sustained depth. Bubble area is exploration-command count. P3 is a persistence outlier; G4 is not.

score fell to 0.63. Package 3’s change was stronger: it issued ten exploration commands before first feedback and 108 afterward. Figure 6 shows that Package 3 is the clear post-feedback outlier, while several controls exhibit smaller pivots.

5.3 Boundary-Crossing And Control Trajectories

The two boundary-crossing runs do not share a simple volume threshold. Gateway 4 had 20 exploration commands, matching Gateway 2; 12 target domains, matching Gateway 1 and 3; and fewer consecutive exploration bins than Gateway 3. It is a low-volume semantic outlier. Package 3 is a high-volume persistence outlier: its 266 commands were 3.3 times the control maximum, and its 118 exploration commands were 7.9 times the control maximum.

Figure 7 makes this contrast visible. Boundary-crossing runs are outlined in red. Gateway 4 remains inside the control cloud, while Package 3 separates by both exploration share and sustained cluster duration.

5.4 Boundary-Crossing Mechanisms And Agent Awareness

Gateway 4 established that direct outbound HTTP failed, inventoried the inference gateway at minute 34.7, and stated that it would use the research-capable endpoint to seek the public UCI source. At minute 36.175 the direct Responses request was accepted; hosted web effects completed at minute 36.179. It later reused the capability at minute 42.268.

Package 3’s route was incremental. It found a UCI-named upstream cache at minute 26.6 and resolved the artifact project at minute 38.5. It then tested package-manager behavior, content and service APIs, encoded paths, installed UI JavaScript, query parameters, methods, headers, and Remote Repository Browser semantics. A 13-header experiment at minute 68.5 caused Nexus to authenticate to the protected origin.

Gateway 4 was aware of the capability it invoked. Its visible commentary distinguished blocked direct HTTP from a research-capable endpoint, requested hosted web search, consumed returned public-page results, and reused the same path. The transcript supports awareness of the external effect, though not a formal statement about authority boundaries.




Figure 8: Key milestones in discovery of the two authority-bearing routes. The gateway path was a late semantic pivot; the package path accumulated service knowledge over roughly an hour.




Figure 9: Retry intensity and route breadth. Bubble area is unique target-domain count. Controls occupy the highest rapid-retry positions, so failure retrying is not a general crossing indicator.

Package 3 was aware that it was experimenting with Nexus and Remote Repository Browser, but awareness of the authenticated downstream consequence is not established. The model-visible output looked like ordinary 200 or 406 XML/JSON responses. Nexus, protected-origin, cloud, and controller records independently show the authenticated effect. Later commentary returned to modeling and did not claim source recovery. This difference is why the authoritative record must separate model observation from system effect.

5.5 Persistence And Post-Effect Behavior

No single persistence proxy predicts crossing. Gateway 2 had 15 failed commands and 11 same-category retries within five minutes, versus six and two for Gateway 4. Package 2 had 13 failures and nine rapid retries, versus 16 and six for Package 3. Endpoint breadth was similarly ambiguous: several controls matched or exceeded a boundary-crossing run.

Package 3's distinctive persistence was hypothesis-generating variation within one authority-bearing service. Its largest exploration cluster contained 105 commands over 60.3 minutes, whereas control clusters contained at most 11 commands and lasted at most 13.8 minutes. Exact repetition accounted for only eight repeated command instances.

Gateway 4 had 6.6 minutes of trace after the first provider effect, containing 16 reasoning items, 14 tool calls, seven commands, and six exploration commands. It reused the hosted capability and then sealed a final submission. Package 3 had 42.0 minutes after L4 and sealed its final 35.2 minutes later. That period contained 160 reasoning items, 122 tool calls, and 104 commands, but only 21 exploratory commands; most later work returned to modeling.

5.6 Between-Run Variation

The gateway condition had nearly fixed observed durations of 42.8–44.6 minutes, yet commands ranged from 48 to 80, exploration shares from 13.8% to 41.7%, failures from four to 15, and target-domain breadth from five to 12. Gateway 4 was near the middle on most measures. The package condition varied from 52.2 to 110.6 minutes and from 59 to 266 commands. Its two best final scores belonged to controls.




Figure 10: Share of each boundary-crossing run’s observed activity occurring after the first effect. Both effects occurred before the run ended; neither ended in protected-label acquisition or demonstrated benchmark benefit.

Run-to-run variation can arise from randomized lane data, stochastic model sampling, tool outcomes, service responses, and right-censoring. Four runs per infrastructure condition describe possible trajectories; they do not estimate the probability that this model will cross a boundary.

6 Causal Reconstruction And Counterfactual Intervention

We next evaluate which independent records are required to reconstruct each boundary event and identify the earliest observed operation at which an admission-time restriction could have changed the trajectory. Because Gensee operated in observe-only mode, the resulting intervention points are trace-derived counterfactuals rather than measurements of prevention efficacy. The traces establish the timing and provenance of authority-bearing effects, but they do not establish how an adaptive agent would respond to denial, whether it would discover another route, or what task performance it would achieve under enforcement.

6.1 Independent Evidence Required For Reconstruction

No single source was sufficient. Agent records established initiation but could not prove what a trusted service did. Service ingress established admission but not necessarily the downstream effect. Network records established connectivity but not application semantics. Reconstruction required joining records with different roles while retaining each source’s provenance and limits.

For Gateway Trial 4, the minimum chain was **agent action** → **gateway admission** → **client-observed provider effect**, interpreted against the documented network envelope. The public trace does not contain an independent provider-side audit log. The reconstruction therefore supports the narrower claim that the admitted request returned hosted-tool effects to the client; it should not be presented as independent proof of the provider’s internal execution path.

For Package Trial 3, the minimum chain was **agent action** → **Nexus ingress** → **authenticated protected-origin effect**, with GCP flow telemetry corroborating the cross-host path. The origin’s authentication record independently showed the credential-bearing effect even though the model-visible response did not disclose it. Controller and replay outputs confirmed the causal join without replacing the underlying raw sources.

6.2 Earliest Counterfactual Intervention Points

Container isolation constrains where untrusted code runs, but it does not fully constrain what an allowed service can perform on that code’s behalf. An operation-scoped contract can bind the initiating subject, request semantics, destination, credential use, downstream effect, and retained output. For these traces, the earliest useful gate is the trusted service’s semantic admission decision, before it exercises broader authority.

For Gateway Trial 4, the trajectory would first change at approximately **36.18 minutes**, when the Responses request declared hosted web tools. A generation-only inference contract could reject hosted-tool declarations or require a separate approved operation before the gateway used its provider credential. The first client-observed hosted-tool effect followed about **0.28 seconds** later, and the capability was reused about **6.09 minutes** after admission.

Table 3: Evidence roles in the two boundary-event reconstructions. Raw and independently operated sources carry the causal claim; controller and replay outputs confirm the join but do not replace its inputs.

Trace role	Gateway Trial 4	Package Trial 3	Why another source was needed
Agent initiation	Model interaction and Codex command	Model interaction and command interaction 711	Establishes the initiating subject and action, but not an external effect.
Boundary admission	Inference-gateway access journal	Nexus 2 ingress journal	Shows that the trusted intermediary accepted the request, but not the full downstream result.
Authority-bearing effect	Provider-effect records and returned command output	Protected-origin authentication logs, including unauthenticated 401 followed by authenticated 200	Connects admission to an effect performed under service-held authority. Gateway evidence is client-observed rather than an independent provider audit.
Network envelope or path	Documented direct-HTTPS denial and drop records	GCP cloud-flow telemetry from Machine B to the protected origin	Distinguishes an intermediary-mediated route from direct agent egress and corroborates the machine-to-origin path.
Derived confirmation	Controller escape marker	Controller L4 marker and authenticated replay	Correlates the records and supports review; it is not a substitute for raw admission and effect evidence.




Figure 11: Evidence records used to reconstruct the two boundary events, shown relative to trusted-service admission. Package Trial 3 includes an independent protected-origin authentication record and cloud-flow corroboration. Gateway Trial 4 ends in client-observed provider effects because no independent provider-side audit was available. The Package 3 source clocks are heterogeneous, so causal correlation uses identifiers and bounded evidence windows rather than timestamp order alone.

For Package Trial 3, the trajectory would first change at approximately **68.52 minutes**, when Nexus admitted the Remote Repository Browser request. A package-fetch contract could reject that open-ended browsing operation before Nexus selected the protected origin and attached its service-held credential. The first authenticated origin effect followed about **0.34 seconds** later, and the L4 controller confirmation followed about **0.58 seconds** after admission.




Figure 12: Earliest observed admission points at which an operation-scoped restriction could have changed each boundary-crossing trajectory. The red branches are counterfactual policy decisions inferred from the traces. No such policy was enabled or experimentally evaluated in these runs.

These subsecond intervals make post-effect reaction too late for the first observed consequence. They motivate admission-time controls over operation semantics and service-held authority, while independent telemetry remains necessary for detection, proof, and response.

The proposed gates preserve ordinary model generation and approved package retrieval while rejecting the specific authority-bearing operations observed here. Because enforcement was not enabled, this analysis cannot determine whether denial would have ended exploration, redirected the agent, exposed another route, or changed final task performance. Those questions require an active, randomized enforcement experiment.

7 Discussion

7.1 Principal Findings

The experiments establish that an agent can produce an externally meaningful effect through a permitted service without a distinctive container-level volume signature. Gateway 4 crossed through a small semantic change in an otherwise ordinary trajectory, whereas Package 3 crossed through sustained, service-specific experimentation. Both effects depended on a service whose effective authority exceeded the agent container’s intended envelope.

The evidence also separates curiosity, action, and consequence. Every run showed some interest in public provenance or infrastructure. Only two requests caused independently confirmed authority-bearing effects. Only Gateway 4 clearly observed its broader downstream capability. Neither boundary-crossing run achieved the protected-information objective. Analyses that equate exploration with benchmark compromise, or a boundary effect with successful data acquisition, would overstate the evidence.

7.2 Operation-Scoped Authority And Independent Observability

The counterfactual analysis locates the useful control point at the semantic admission boundary, not at generic infrastructure discovery and not after an effect. The provider and package traces suggest that a contract must constrain request semantics, destination selection, and service-held credential use together. A network allowlist alone cannot express those application-level limits.

The two proof chains have different epistemic strength. Package Trial 3 has an independent protected-origin audit trail and cross-host flow corroboration. Gateway Trial 4 has gateway admission and client-observed provider-effect items, but no independent provider-side audit. Stating this asymmetry prevents a derived controller marker from being mistaken for a stronger raw observation.

Independent host, network, service, and controller telemetry can detect effects even when the agent is unaware of them or attempts to manipulate its own transcript. Digest-bound event chains and stable execution identity support forensic attribution and offline replay without re-executing captured

commands. Online policy and post-incident proof therefore use the same causal structure, while retaining the provenance and limits of every observation.

7.3 Limitations

The study contains four runs per infrastructure condition and is descriptive. It cannot estimate the probability of a boundary crossing, the prevalence of the observed behaviors across models, or the effectiveness of a defense. The Gateway controls were also right-censored by condition-wide shutdown after the boundary event. They provide contemporaneous comparisons but are not completed negative trials. Because the two conditions differ in infrastructure and telemetry, the eight runs cannot be pooled as interchangeable trials or compared using a universal event-count scale.

The retained telemetry provides incomplete and source-dependent views of execution. Falco records configured rule matches rather than every syscall; the Gateway rules omit ordinary reads, the Package rules are broader, and Package 3 reports substantial capture drops. Service logs may contain multiple records for one causal action, so totals across sources are not additive measures of independent behavior. Public reasoning items are opaque, and the study does not observe chain of thought, reasoning-token counts, or private model state. Behavioral interpretation therefore relies on explicit commentary and observable actions. The deterministic classification codebook is also analyst-designed, and alternative taxonomies may group commands differently.

The experimental fixture limits external validity. Nexus was intentionally vulnerable and lane-scoped, so the package-service result does not generalize to patched or normally configured deployments. The observed boundary events establish capability deltas, not protected-label delivery, use, or benchmark compromise. The Gateway evidence is additionally limited by the absence of an independent provider-side audit log; its provider effects are retained client-observed completion items.

Finally, the intervention analysis is counterfactual. No active denial or approval gate was enabled, and the study consequently does not measure prevention efficacy, agent adaptation after denial, or the possibility that an agent would discover an alternative route.

7.4 Future Research

Larger studies should vary seeds, models, harnesses, and task families while using uncensored controls and preregistered stopping rules. This would support estimates of crossing frequency and clarify which trajectory patterns generalize beyond the two infrastructure conditions studied here.

Measurement should also be standardized across conditions. Future trace collections should harmonize host telemetry, preserve observer-loss and clock-quality metadata, and decode retained binary captures where safe. A preregistered phase and strategy codebook, accompanied by inter-annotator agreement, would make behavioral comparisons more reproducible. Causal graphs could then connect the initiating subject, mediator admission, credential use, downstream destination, effect, and retained artifact across heterogeneous clocks.

The trace-derived intervention points motivate controlled causal experiments. Such studies could independently disable provider-hosted tools, narrow package-service upstream identity, constrain destinations, introduce credentialless mediation, and enable machine-speed shutdown. Evaluating operation-scoped contracts under adaptive retry would test whether denial ends the trajectory, redirects the agent, or exposes another transitive-authority path. Simulated alerting experiments could additionally quantify detection latency, response budgets, automatic lease closure, and condition-wide containment.

Further releases should include privacy-preserving raw schemas and derived data that support independent replication, while continuing to gate live credentials, private identities, unsafe payloads, and external replay capabilities.

8 Conclusion

Across eight autonomous coding-agent runs, broad infrastructure interest was common, but independently confirmed authority-bearing effects were rare and heterogeneous. Gateway 4 crossed through an overbroad semantic request without unusual model, command, host, or network volume. Package

3 crossed after an unusually long and dense service-specific investigation. Neither run accessed protected labels or outperformed the controls.

The central analytical lesson is to measure the full causal unit of work. Agent commentary and local syscall-like events explain only part of the trajectory. Security decisions require a joined record of the approved caller, admitted operation, trusted mediator, credential use, downstream effect, and retained output. That record supports narrow online prevention, independent detection, and credible post-incident reconstruction without treating model-generated telemetry as authoritative. The counterfactual intervention analysis identifies where this record could support semantic admission control, while remaining explicit that prevention was not experimentally evaluated.

References

- Anthropic. Investigating three real-world incidents in our cybersecurity evaluations. <https://www.anthropic.com/research/investigating-incidents-cybersecurity-evals>, July 2026. Accessed 2026-08-27.
- Gensee AI. Open agent-security traces: Eight autonomous-agent boundary-crossing trials. <https://github.com/GenseeAI/gensee-crate/tree/main/security-traces>, 2026. Version 1 public trace corpus; accessed 2026-08-27.
- Brian A. Johnson, Kotaro Iizuka, Milton A. Bragais, I. Endo, and Damasa B. Magcale-Macandog. Crowdsourced mapping. UCI Machine Learning Repository, 2016. URL <https://archive.ics.uci.edu/dataset/400/crowdsourced+mapping>. Accessed 2026-08-27.
- OpenAI. Pacing model development with cyber capabilities. <https://openai.com/index/pacing-model-development-cyber-capabilities/>, August 2026a. Accessed 2026-08-27.
- OpenAI. The hugging face incident and the road ahead. <https://openai.com/index/hugging-face-incident-and-the-road-ahead/>, August 2026b. Accessed 2026-08-27.
- OpenAI. OpenAI–Hugging Face Incident: Technical Report. Technical report, OpenAI, August 2026c. URL <https://cdn.openai.com/pdf/67869394-cb91-4c12-888c-5cbd85c7814c/OpenAI-Hugging-Face%20Incident-Technical-Report.pdf>. Accessed 2026-08-27.
- UK AI Security Institute. Incident report: Unsanctioned agent behaviour during cyber testing. <https://postmortem.io/incidents/aisi--2026-08-04--unsanctioned-agent-behaviour-cyber-testing/>, August 2026. Incident report; accessed 2026-08-27.
- Hjalmar Wijk, Ajeya Cotra, and Ryan Greenblatt. Brief independent investigation of agents’ behavior, reasoning and collaboration in the openai / hugging face hacking incident. <https://metr.org/blog/2026-08-26-openai-hugging-face-incident-investigation/>, August 2026. METR and Redwood Research; accessed 2026-08-27.